

CS 3400 Machine Learning - Lab 4: Feedforward Neural Network

Stuart Harley

Introduction

Neural Networks are a type of machine learning model that were inspired by the human brain. The networks are organized into layers and within these layers are neurons. Each neuron implements a linear model whose output is passed through an activation function. A multilayer perceptron is a class of feedforward neural network. In this lab, we will be experimenting with MLP's to classify the iris dataset.

Reflection Questions


Problem 1

1. The parameters to the MLPClassifier that we used in this lab are as follows. `hidden_layer_sizes` : The size (number of neurons) of the hidden layers within the perceptron. `max_iter` : The maximum number of iterations used during the training of this network. `solver` : The formula used for weight optimization. In this lab we used 'lbfgs' which is an optimizer in the family of quasi-Newton methods.

```
In [1]: from IPython.display import Image
print('2.')
Image('MLP Drawing.jpeg', width=400, height=400)
```

2.

Out[1]:



1. The activation function used for the nodes is the relu function.

Problem 2

1. `mlp_petals.coefs[0]` is a 2×4 matrix. `mlppetals.intercepts[0]` is a 4×1 vector. These dimensions come from the size of the MLP. There are 2 coefficients (B_1 and B_2) for each neuron in the hidden layer so it is 2×4 . There is 1 intercept for each neuron in the hidden layer so it is 4×1 .
1. `mlp_petals_models` is a 4×3 matrix. The dimensions correspond to combining the previously mentioned matrices. When you stack a 4×2 and a 4×1 matrix, you end up with a 4×3 matrix.
1. Separate Setosa from the rest : Planes 1 or 3 from the petal data.

Separate Versicolor from the rest : Planes 2 and 3 from the petal data.

Separate Virginica from the rest : Plane 3 from the petal data.

Problem 3

1. A relu function modifies the output by setting any output < 0 to 0. Any output > 0 stays the same. This impacts the interpretation because anything on one side of the plane by the relu function becomes on the plane.

1. Based on the heatmap plots.

Separate Setosa from the rest : Neuron 1 and 4

Separate Versicolor from the rest : Neuron 3

Separate Virginica from the rest : Neuron 2

1. The ReLu activation layer makes it easier to use decision boundaries because you don't have to worry about how far on the negative side of the plane the value is, only the positive, because all negative values become 0. However, if you have a point directly on the plane, it could make things more difficult because that 0 value would be the same as some large negative value that has become 0.

Problem 4

1. The accuracy of the 4 features is 92% while the accuracy of the transformed features is 97.33%. Therefore, the transformed features produced a more accurate model.

Importing Libraries

```
In [2]: from sklearn import datasets
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import confusion_matrix
from neurons import *
```

1)

Loading Iris data set

```
In [3]: iris = datasets.load_iris()
```

Training a Multilayer Perceptron on the petal lengths and widths

```
In [4]: scaled_X = StandardScaler().fit_transform(iris.data)
mlp_petals = MLPClassifier(hidden_layer_sizes=(4,), max_iter=1000, solver='lbfgs')
mlp_petals.fit(scaled_X[:, 2:], iris.target)
```

```
Out[4]: MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
beta_2=0.999, early_stopping=False, epsilon=1e-08,
hidden_layer_sizes=(4,), learning_rate='constant',
learning_rate_init=0.001, max_fun=15000, max_iter=1000,
momentum=0.9, n_iter_no_change=10, nesterovs_momentum=True,
power_t=0.5, random_state=None, shuffle=True, solver='lbfgs',
tol=0.0001, validation_fraction=0.1, verbose=False,
warm_start=False)
```

Creating a second Multilayer Perceptron for the sepal lengths and widths

```
In [5]: mlp_sepals = MLPClassifier(hidden_layer_sizes=(4,), max_iter=1000, solver='lbfgs')
mlp_sepals.fit(scaled_X[:, 0:2], iris.target)
```

```
Out[5]: MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
beta_2=0.999, early_stopping=False, epsilon=1e-08,
hidden_layer_sizes=(4,), learning_rate='constant',
learning_rate_init=0.001, max_fun=15000, max_iter=1000,
momentum=0.9, n_iter_no_change=10, nesterovs_momentum=True,
power_t=0.5, random_state=None, shuffle=True, solver='lbfgs',
tol=0.0001, validation_fraction=0.1, verbose=False,
warm_start=False)
```

2)

Extracting the weight vectors for the hidden layers

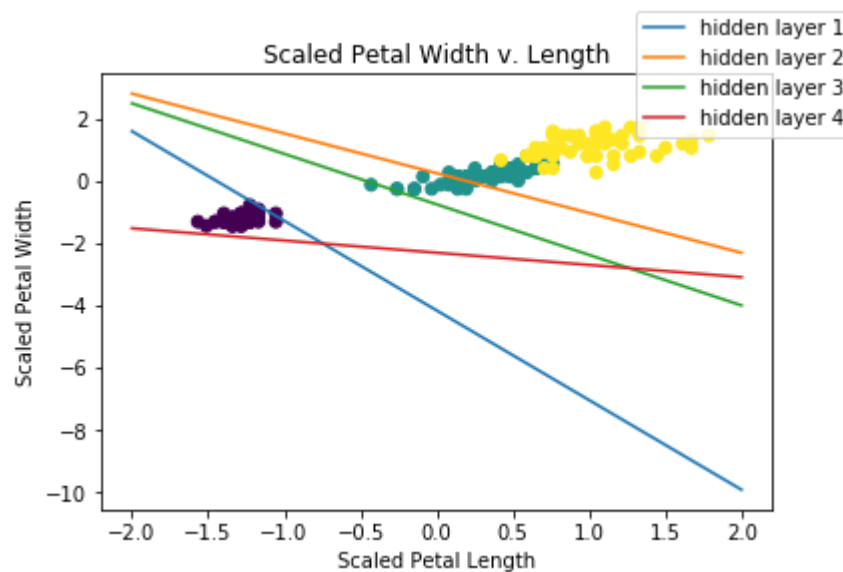
```
In [6]: mlp_petals_models = np.vstack([mlp_petals.intercepts_[0], mlp_petals.coefs_[0]]).T
mlp_sepals_models = np.vstack([mlp_sepals.intercepts_[0], mlp_sepals.coefs_[0]]).T
```

The equation representing the hidden layers: $0 = B_0 + B_1 \cdot x_1 + B_2 \cdot x_2$ solved for x_2 becomes $x_2 = (-B_0 - B_1 \cdot x_1) / B_2$

```
In [7]: def x2(B0, B1, B2, x1):
return (-B0 - B1 * x1) / B2
```

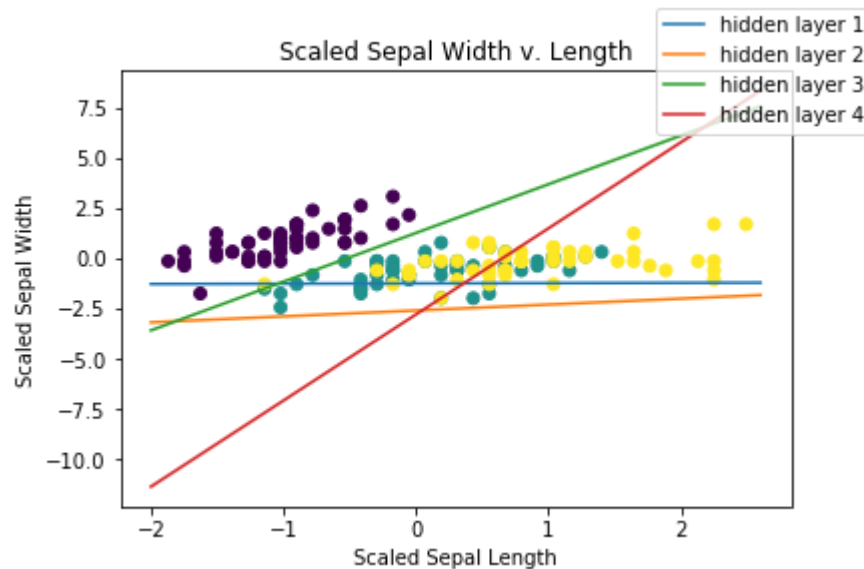
Plotting scaled petal features along with the planes for the 4 hidden layers.

```
In [8]: x1 = np.linspace(-2, 2, 100)
fig, axes = plt.subplots()
axes.scatter(scaled_X[:,2], scaled_X[:,3], c=iris.target)
axes.plot(x1, x2(mlp_petals_models[0,0], mlp_petals_models[0,1], mlp_petals_models[0,2], x1), label='hidden layer 1')
axes.plot(x1, x2(mlp_petals_models[1,0], mlp_petals_models[1,1], mlp_petals_models[1,2], x1), label='hidden layer 2')
axes.plot(x1, x2(mlp_petals_models[2,0], mlp_petals_models[2,1], mlp_petals_models[2,2], x1), label='hidden layer 3')
axes.plot(x1, x2(mlp_petals_models[3,0], mlp_petals_models[3,1], mlp_petals_models[3,2], x1), label='hidden layer 4')
axes.set_xlabel('Scaled Petal Length')
axes.set_ylabel('Scaled Petal Width')
axes.set_title('Scaled Petal Width v. Length')
fig.legend();
```



Plotting scaled sepal features along with the planes for the 4 hidden layers.


```
In [9]: x1 = np.linspace(-2, 2.6, 100)
fig, axes = plt.subplots()
axes.scatter(scaled_X[:,0], scaled_X[:,1], c=iris.target)
axes.plot(x1, x2(mlp_sepals_models[0,0], mlp_sepals_models[0,1], mlp_sepals_models[0,2], x1), label='hidden layer 1')
axes.plot(x1, x2(mlp_sepals_models[1,0], mlp_sepals_models[1,1], mlp_sepals_models[1,2], x1), label='hidden layer 2')
axes.plot(x1, x2(mlp_sepals_models[2,0], mlp_sepals_models[2,1], mlp_sepals_models[2,2], x1), label='hidden layer 3')
axes.plot(x1, x2(mlp_sepals_models[3,0], mlp_sepals_models[3,1], mlp_sepals_models[3,2], x1), label='hidden layer 4')
axes.set_xlabel('Scaled Sepal Length')
axes.set_ylabel('Scaled Sepal Width')
axes.set_title('Scaled Sepal Width v. Length')
fig.legend();
```



3)

Creating a mesh grid from (-2, 2) and plotting it as a scatter plot

```
In [10]: x = np.linspace(-2, 2, 100)
xv, yv = np.meshgrid(x, x)
plt.scatter(xv, yv, s=1, c='k');
```




Using the Input and Neuron classes from neruons.py to calculate the value for the first hidden layer neuron at each grid point (X).

```
In [11]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_petals_models[0, :])
pred = p_layer.predict(X)
```

Plotting the model outputs as a heatmap.


```
In [12]: pred = pred.reshape(100, 100)
heatmap = sns.heatmap(pred)
```



Repeating for the 3 remaining neurons in the hidden layer.


Second Neuron.

```
In [13]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_petals_models[1, :])
pred = p_layer.predict(X).reshape(100,100)
heatmap = sns.heatmap(pred)
```



Third Neuron

```
In [14]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_petals_models[2, :])
pred = p_layer.predict(X).reshape(100, 100)
heatmap = sns.heatmap(pred)
```



Fourth Neuron



```
In [15]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_petals_models[3, :])
pred = p_layer.predict(X).reshape(100, 100)
heatmap = sns.heatmap(pred)
```



Repeating for the hidden layers of the Sepals model

First Neuron

```
In [16]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_sepals_models[0, :])
pred = p_layer.predict(X).reshape(100, 100)
heatmap = sns.heatmap(pred)
```



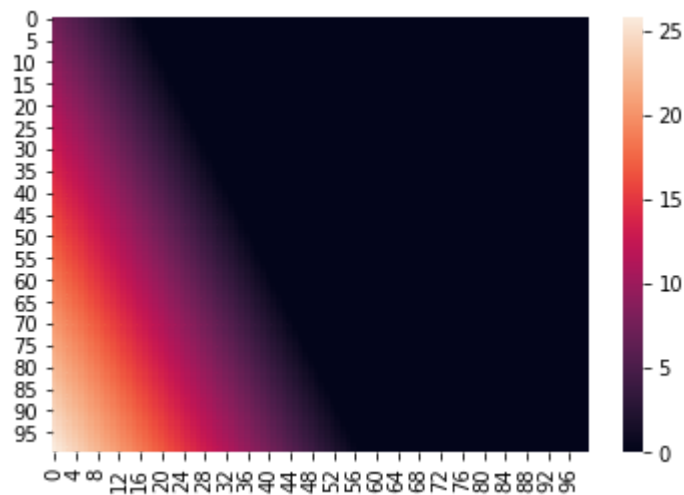
Second Neuron

```
In [17]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_sepals_models[1, :])
pred = p_layer.predict(X).reshape(100, 100)
heatmap = sns.heatmap(pred)
```




Third Neuron

```
In [18]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_sepals_models[2, :])
pred = p_layer.predict(X).reshape(100, 100)
heatmap = sns.heatmap(pred)
```



Fourth Neuron

```
In [19]: X = np.hstack([xv.reshape(-1, 1), yv.reshape(-1, 1)])
input = Input()
p_layer = Neuron([input], mlp_sepals_models[3, :])
pred = p_layer.predict(X).reshape(100, 100)
heatmap = sns.heatmap(pred)
```



4)

Using the Input, Neuron, and HStack classes with the weights from the MLP model to recreate the hidden layer.

```
In [20]: input = Input()
p_layer_1 = Neuron([input], mlp_petals_models[0, :])
p_layer_2 = Neuron([input], mlp_petals_models[1, :])
p_layer_3 = Neuron([input], mlp_petals_models[2, :])
p_layer_4 = Neuron([input], mlp_petals_models[3, :])
stacked = HStack([p_layer_1, p_layer_2, p_layer_3, p_layer_4])
```

Predicting the transformed values to create a transformed feature matrix.

```
In [21]: transformed_petals_X = stacked.predict(scaled_X[:, 2:])
```

Repeating for the sepal MLP model

```
In [22]: input = Input()
p_layer_1 = Neuron([input], mlp_petals_models[0, :])
p_layer_2 = Neuron([input], mlp_petals_models[1, :])
p_layer_3 = Neuron([input], mlp_petals_models[2, :])
p_layer_4 = Neuron([input], mlp_petals_models[3, :])
stacked = HStack([p_layer_1, p_layer_2, p_layer_3, p_layer_4])
transformed_sepals_X = stacked.predict(scaled_X[:, 0:2])
```

Combining the two transformed feature matrices into a new feature matrix with 8 columns.

```
In [23]: transformed_stacked = np.hstack((transformed_petals_X, transformed_sepals_X))
```

Training two logistic regression models. One on the original 4 features and one on the new transformed feature matrix with 8 columns.

```
In [24]: four_feat = SGDClassifier(loss='log')
four_feat.fit(scaled_X, iris.target)
```

```
Out[24]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                        early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
                        l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=1000,
                        n_iter_no_change=5, n_jobs=None, penalty='l2', power_t=0.5,
                        random_state=None, shuffle=True, tol=0.001,
                        validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [25]: eight_feat = SGDClassifier(loss='log')
eight_feat.fit(transformed_stacked, iris.target)
```

```
Out[25]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                        early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
                        l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=1000,
                        n_iter_no_change=5, n_jobs=None, penalty='l2', power_t=0.5,
                        random_state=None, shuffle=True, tol=0.001,
                        validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the two models using accuracy and confusion matrices.

```

In [26]: def plot_confusion_matrix(cm, target_names, title='Confusion matrix',
                                     cmap=None, normalize=True):
    """
    given a sklearn confusion matrix (cm), make a nice plot

    Citation
    -----
    http://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_confusion\_matrix.html

    """
    import itertools

    accuracy = np.trace(cm) / np.sum(cm).astype('float')
    misclass = 1 - accuracy

    if cmap is None:
        cmap = plt.get_cmap('Blues')

    plt.figure(figsize=(6, 4))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()

    if target_names is not None:
        tick_marks = np.arange(len(target_names))
        plt.xticks(tick_marks, target_names, rotation=45)
        plt.yticks(tick_marks, target_names)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 1.5 if normalize else cm.max() / 2
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        if normalize:
            plt.text(j, i, "{:0.4f}".format(cm[i, j]),
                     horizontalalignment="center",
                     color="white" if cm[i, j] > thresh else "black")
        else:
            plt.text(j, i, "{:,}".format(cm[i, j]),
                     horizontalalignment="center",
                     color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label\naccuracy={:0.4f}'.format(accuracy))
    plt.show()

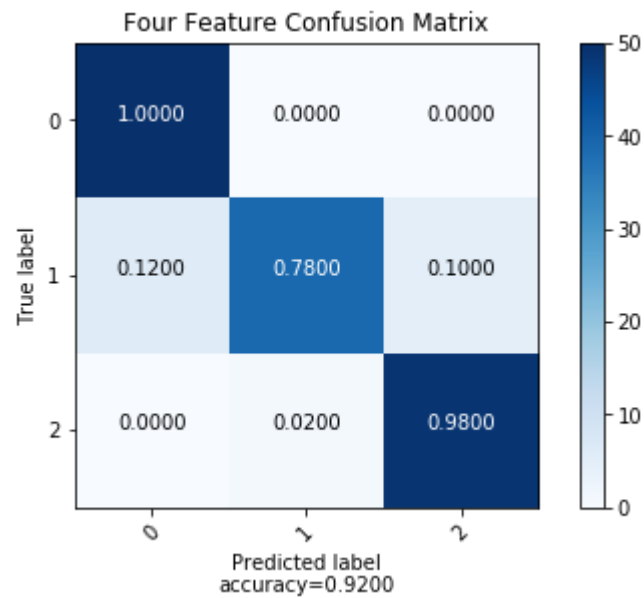
```

```

In [27]: four_confusion = confusion_matrix(iris.target, four_feat.predict(scaled_X))
eight_confusion = confusion_matrix(iris.target, eight_feat.predict(transformed_
_stacked))

```

```
In [28]: plot_confusion_matrix(four_confusion, [0, 1, 2], title='Four Feature Confusion Matrix')
```



```
In [29]: plot_confusion_matrix(eight_confusion, [0, 1, 2], title='Eight Feature Confusion Matrix')
```

